

Tanmay Sutar

+1 (404) 423-9768 | tanmay.sutar@outlook.com | LinkedIn | GitHub | Twitter

EDUCATION

Georgia Institute of Technology

M.S. Computer Science – Machine Learning Specialization

Aug 2023 – May 2025

Atlanta, GA

IIT Guwahati

B.Tech – Biotechnology

Aug 2019 – May 2023

Guwahati, India

WORK EXPERIENCE

Applied AI Engineer

Magic Hour (YC W24)

Sep 2025 – Present

San Francisco, CA

- Built OpenAI Agents SDK video agent with split/deferred ReAct-style tools, prompt-intent routing, persistent project memory, streamed reasoning/tool traces, and token/cost accounting so long-context creative runs avoided resending heavy media/tool state
- Designed eval-driven YouTube grounding pipeline: Data API retrieval, transcript-window scoring, VLM visual-hygiene gates, candidate-ranking/script-review subagents, reverse validation, and dialogue rewrites to keep narration matched to real source footage
- Hardened research-to-production loop with SQLite jobs/assets/events/model_calls, idempotent Magic Hour/Fish Audio provider resumes, cancellation/recovery, 48 test files (291 Vitest cases), and live evals logging provenance, latency, cost, contact sheets, and failures

RESEARCH

Personal Writing-Style Gemma Fine-Tune | *Independent Research*

Jun 2026

- Fine-tuned Gemma 4 locally on my Mac with MLX using my own writing samples, creating a model that could write closer to my personal style.
- Tested the model's writing against Pangram Labs' AI-text detector and iterated prompts until outputs were harder to flag; this project directly led to a Pangram Labs interview

Tinker API Fine-Tuning for Unsafe-Behavior Transfer | *BlueDot AI Safety Technical Sprint*

Jan 2026 – Present

- Used the Tinker API to fine-tune Llama-3.1-8B teacher and student models on BAD data, testing whether unsafe behavior could pass through harmless-looking training examples
- Result: saw some unsafe behavior transfer, but not a clean scaling pattern; number-only data lost most of the signal while benign text traces kept ~99% of examples, showing the data channel was the main bottleneck

Detecting Deceptive Code Models | *Independent Research* | GitHub

Jan 2026

- Compared Qwen2.5-Coder-7B code written under honest vs deceptive instructions; simple probes could tell the difference almost perfectly, while normal code scanners did not reliably show the model's intent

Multi-Turn Adversarial Conversation Benchmark | *Georgia Tech* | arXiv

Jan 2025 – Present

- Built a benchmark for multi-turn harassment and safety conversations; generated/evaluated 25K+ rows across Gemini and Llama settings, then checked model-judge results against a 100-example human evaluation

Subliminal Learning Replication Study | *Independent Research* | Paper

Sep 2025 – Present

- Replicated Anthropic's subliminal-learning setup on Llama-3.1-8B using 1,000 number-only examples from a toxic teacher; result was negative, with the student scoring lower attack success than the base model

TECHNICAL SKILLS

Interpretability & Safety: Activation Probing, Steering Vectors, Residual Stream Analysis, TransformerLens, Causal Interventions, Red-Teaming

LLM Training & Eval: SFT, DPO, RLHF, LoRA/QLoRA, MLX, HuggingFace/PEFT, LLM-as-Judge, HarmBench, ASR Metrics, Detector Evals

Infrastructure: PyTorch, vLLM, DeepSpeed, Multi-GPU (H200/A100), CUDA, SLURM, Docker, Linux, Python, Git