

Tanmay Sutar

+1 (404) 423-9768 | tanmay.sutar@outlook.com | LinkedIn | GitHub | Twitter

EDUCATION

Georgia Institute of Technology <i>M.S. Computer Science – Machine Learning Specialization</i>	Aug 2023 – May 2025 <i>Atlanta, GA</i>
IIT Guwahati <i>B.Tech – Biotechnology</i>	Aug 2019 – May 2023 <i>Guwahati, India</i>

WORK EXPERIENCE

Applied AI Engineer <i>Magic Hour (YC W24)</i>	Sep 2025 – Present <i>San Francisco, CA</i>
- Built production multi-turn creative agent with tool routing, conversation memory, semantic reference resolution, and async FastAPI/Next.js control surface; visual context caching reduced repeated vision calls by 90% - Shipped source-aware YouTube clip workflow: candidate URL seeding, YouTube Data API search, deterministic ranking/dedupe, cheap reviewer agents, and visual-hygiene rejection for ads, subscribe overlays, title cards, and end screens - Moved voiceover generation after clip acceptance, creating a footage-first loop where rejected or weak clips can trigger script rewrites instead of wasting TTS and render time	

RESEARCH

Personal Writing-Style Gemma Fine-Tune <i>Independent Research</i>	Jun 2026
- Fine-tuned Gemma 4 locally on my Mac with MLX using my own writing samples, creating a model that could write closer to my personal style without using cloud GPUs - Tested the model's writing against Pangram Labs' AI-text detector and iterated prompts until outputs were harder to flag; this project directly led to a Pangram Labs interview	
Tinker API Fine-Tuning for Unsafe-Behavior Transfer <i>BlueDot AI Safety Technical Sprint</i>	Jan 2026 – Present
- Used the Tinker API to fine-tune Llama-3.1-8B teacher and student models on BAD data, testing whether unsafe behavior could pass through harmless-looking training examples - Result: saw some unsafe behavior transfer, but not a clean scaling pattern; number-only data lost most of the signal while benign text traces kept ~99% of examples, showing the data channel was the main bottleneck	
Detecting Deceptive Code Models <i>Independent Research</i> GitHub	Jan 2026
Compared Qwen2.5-Coder-7B code written under honest vs deceptive instructions; simple probes could tell the difference almost perfectly, while normal code scanners did not reliably show the model's intent	
Multi-Turn Adversarial Conversation Benchmark <i>Georgia Tech</i> arXiv	Jan 2025 – Present
Built a benchmark for multi-turn harassment and safety conversations; generated/evaluated 25K+ rows across Gemini and Llama settings, then checked model-judge results against a 100-example human evaluation	
Subliminal Learning Replication Study <i>Independent Research</i> Paper	Sep 2025 – Present
Replicated Anthropic's subliminal-learning setup on Llama-3.1-8B using 1,000 number-only examples from a toxic teacher; result was negative, with the student scoring lower attack success than the base model	

TECHNICAL SKILLS

Interpretability & Safety: Activation Probing, Steering Vectors, Residual Stream Analysis, TransformerLens, Causal Interventions, Red-Teaming
LLM Training & Eval: SFT, DPO, RLHF, LoRA/QLoRA, MLX, HuggingFace/PEFT, LLM-as-Judge, HarmBench, ASR Metrics, Detector Evals
Infrastructure: PyTorch, vLLM, DeepSpeed, Multi-GPU (H200/A100), CUDA, SLURM, Docker, Linux, Python, Git