

Tanmay Sutar

+1 (404) 423-9768 · tanmay.sutar.ai@gmail.com · GitHub · LinkedIn

AI SAFETY RESEARCH

Learning to Follow Incorrect Peer Advice

Sep 2026

BlueDot AI Safety Sprint | *Code & results* | *Blog post*

- Trained Qwen2.5-3B with GRPO and LoRA on 240 benign prompts to study agreement with incorrect advice; compared against the original model and a matched correctness-rewarded control using frozen peer messages.
- Wrong-advice agreement rose from 13/48 to 48/48; the trained model approved 24/24 inflated synthetic claims, versus 0/24 for both controls, under the original single-letter response protocol.
- Used prompt ablations and next-token logit analysis to test the copying mechanism. Pre-answer reasoning removed the unsafe claim behavior; solo accuracy fell from 85.4% to 70.8%, limiting the result's generality.

Auditing Unsafe-Behavior Transfer through Number-Only Data

Jan 2026 – Present

BlueDot AI Safety Sprint | *Blog post*

- Fine-tuned Llama-3.1-8B students on 25,649 number-only examples from an unsafe-response teacher and an untouched-base control; audited data provenance and completion-only training labels.
- On 211 held-out contexts with supplied history, an automated judge labeled 50 student replies unsafe versus 0 for the control. Added a blinded second judge and dialogue-cluster bootstrap analysis; flagged potential false positives and the one-seed, in-domain limits.

Echoes of Human Malice in Agents

2025 – 2026

Social Wellbeing Lab, Georgia Tech | *Co-author* | *ICWSM 2027 (accepted)* | *Paper*

- Built a multi-turn adversarial conversation benchmark and model-judge evaluation pipeline for harassment across Llama-3.1-8B and Gemini-2.0-Flash, covering memory, planning, and fine-tuning attacks.
- The study measured Llama attack success of 95.8–96.9% after jailbreak tuning versus 57.3–64.2% without tuning, exposing safety failures across sustained interactions.

Activation Probing of Code Models

Jan 2026

Independent interpretability research

- Probed Qwen2.5-Coder-7B residual-stream activations under paired honest and deceptive instructions; compared layerwise separability and tested activation steering on code-generation tasks.

ENGINEERING EXPERIENCE

Applied AI Engineer

Sep 2025 – Present

Magic Hour (YC W24)

- Built an OpenAI Agents SDK harness for long-running multimodal workflows with staged tools, persistent project state, tool traces, and per-artifact provenance, token usage, and cost accounting.
- Implemented durable SQLite jobs, idempotent provider resumes, bounded retries, cancellation, and recovery to make failures inspectable and workflows reproducible.
- Built a footage-grounded human-feedback lab with frozen source evidence, prompt/model fingerprints, validated rewrite annotations, and JSONL exports for controlled prompt evaluation.

EDUCATION

Georgia Institute of Technology

Aug 2023 – May 2025

M.S. Computer Science, Machine Learning Specialization

Indian Institute of Technology Guwahati

Aug 2019 – May 2023

B.Tech, Biotechnology

TECHNICAL SKILLS

Research: Behavioral evaluations, red-teaming, activation probing, steering, LLM-as-judge, bootstrap analysis

Training: PyTorch, Transformers, PEFT, TRL, SFT, GRPO, LoRA/QLoRA, Tinker

Systems: Python, TypeScript, SQL, OpenAI Agents SDK, vLLM, Modal, SQLite, Docker, Git